

A

ASSAY · TOKEN TRUST RATINGS

The Independent Rating Agency for Tokens

One letter grade for the gap between what a crypto project promises and what its blockchain proves — with the receipts, over time.

Methodology, market, architecture, business model, legal structure, name, and the road from first commit to end-state.

WORKING NAME

Assay (fallback: TokenFacts) — see Part VII

WEDGE

Base / EVM tokens that make verifiable claims

ENTITY

Wyoming LLC (recommended) — see Part VI

PRODUCT

Algorithmic token trust grades + behavioral alerts

DIFFERENTIATORS

Docs-vs-chain claims engine · public trust time-series

YEAR-1 CASH NEED

≈ \$2,000–4,000 all-in (infra + legal + domains)

CONTENTS

01	Part I — The Product	3
02	Part II — The Methodology	5
03	Part III — Market & Positioning	9
04	Part IV — Technical Architecture & Build Plan	12
05	Part V — Business Model & Financials	15
06	Part VI — Legal, Entity & the ATA Conflict	18
07	Part VII — Name & Domains: What to Buy	21
08	Part VIII — Roadmap to End-State	23
09	Part IX — New Repository Bootstrap & Decision Log	25

The Product

What it is, why the moment is right, and the one sentence the whole business hangs on.

The product is a public trust rating for crypto tokens. Anyone can look up a token and see one legible letter grade (A+ to F), a per-dimension breakdown, and — this is the differentiator — a claim-by-claim comparison of what the project *says* in its website and documents against what its blockchain *actually shows*. Grades update automatically as on-chain reality changes, so every token accumulates a public track record: promises kept, promises broken, admin keys moved, liquidity locks renewed or quietly expired.

THE SENTENCE THE GRADE HANGS ON

The grade measures whether a token can betray you in ways it hasn't disclosed — not whether it is a good investment. Disclosed, purpose-consistent centralization is not a lie (USDC is pausable by design and says so). A memecoin that promises nothing and hides nothing can outscore a "serious" project that lies. And trust needs time: no token — including the founder's own — can score an A on day one.

Why this, why now

The token-safety market is crowded at the bottom and empty exactly where this product sits. Free checkbox scanners (GoPlus, Token Sniffer, RugCheck) tell you whether a contract is a honeypot; they are commoditized to zero. The one large paid player, CertiK, is reputationally compromised because projects pay it for their own scores — the 2008 credit-rating conflict re-run. Research across the July 2026 landscape found **two niches that are verifiably empty**:

1. **LLM-powered docs-vs-chain claim verification.** Nobody systematically extracts "liquidity locked / team vested 4 years / no mint function / treasury is a 3-of-5 multisig" from a project's own marketing and grades the delta against chain state. The closest single feature (Tokenomist's unlock-claims page) covers one claim type on large caps only, with no grade and no alerts.
2. **A public behavioral trust time-series.** Enterprise firms (Hexagate — acquired by Chainalysis for ~\$60M — and Hypernative) monitor admin-key changes privately for paying protocols. No one publishes "this token's trust trajectory over 18 months" as a public, per-token record.

Both gaps are the same product seen from two angles: *verify the promises, then keep watching them*. A time-series moat compounds — every month it runs is a month no competitor can backfill honestly.

Where it comes from

The methodology is the ATA Coin playbook generalized. ATA was engineered so that every trust promise is enforced by code and checkable by anyone — no mint, atomic distribution, on-chain vesting, published policies, an MCP server that answers from chain state. The rating product asks of every other token the questions ATA was built to answer, and the ATA repository contributes directly reusable chain-reading code (Part IV). The affiliation is disclosed, the brands and entities stay separate, and ATA itself is handled under a strict conflict-of-interest policy (Part VI §5).

What the user sees

SURFACE	CONTENT	PRICE
Public site (free, forever)	Letter grade + Trust Model badge, per-dimension breakdown, every claim with its verdict and on-chain evidence link, full methodology, historical grade timeline	\$0
Mention bots (X, Farcaster, Telegram)	"@bot grade \$TOKEN" → grade card + worst finding; the growth loop	\$0
Pro alerts	Real-time: admin key changed, LP lock expiring, vesting wallet moved, docs changed	≈ \$24/mo
API (free tier with attribution; paid tiers)	Machine-readable grades for wallets, screeners, and AI agents (MCP + x402)	\$0 / \$99 / \$399+/mo
Deep-dive reports	Human-reviewed trust-gap due diligence for funds, exchanges, DAOs — buyer-pays, never issuer-pays	≈ \$3,500

THE CREDIBILITY POSTURE (LOAD-BEARING, NON-NEGOTIABLE)

Scores are computed by a published, versioned methodology. **No issuer ever pays for a score. The service never launches a token. The founder never trades around rating changes.** Open methodology + free grades + proprietary pipeline — the L2Beat playbook with a business model attached.

The Methodology

Eight dimensions, a letter scale with hard caps, resistance to gaming, and the proof that the scale is honest: the founder's own token can't get an A.

2.1 The eight dimensions

Each dimension scores 0–100 against published anchors. Computability tags: [CHAIN] = automated from public RPC/explorer data, [LLM] = needs the claims engine, [MANUAL] = human review (the paid tier).

#	DIMENSION	WEIGHT	QUESTION IT ANSWERS	MAINLY
1	Supply Integrity	20	Can the number of tokens change outside disclosed rules? (mint paths, rebasing, hidden fees, latent proxy-mint)	CHAIN
2	Admin Surface & Upgradeability	20	Who can do what to the contract, and what gates them? (owner, roles, proxies, pause/blacklist/fee switches)	CHAIN
3	Disclosure Integrity	15	The docs-vs-chain gap — the differentiator (claims extracted, each verified against chain state)	LLM+CHAIN
4	Insider Float & Vesting	15	How much can insiders sell, and is the lock code or a promise?	CHAIN+LLM
5	Liquidity Permanence	15	Can the pool be pulled? (% burned/locked, locker identity, expiry, depth)	CHAIN
6	Operational Track Record	10	Trust as a time series: age, unannounced admin actions, commitments met	CHAIN+LLM
7	Holder Concentration	5	Top-10/50 share excluding labeled non-circulating addresses; sniper detection	CHAIN
8	Governance Reality	5*	If governance is claimed, is it real or theater? (*only when claimed; else weight returns to #2)	CHAIN+LLM

Supply and admin carry the most weight because they are the two vectors that can take 100% of holder value in one transaction. Concentration gets the least because labeling is the noisiest data — its extremes are handled by hard caps instead.

2.2 The claims engine (dimension 3, the moat)

Pipeline: crawl the project's site, whitepaper, docs, and pinned socials → an LLM extracts claims into a typed schema (claim text, exact quote, source URL, predicate type) → each claim routes to a deterministic verifier → verdict: **VERIFIED / FALSE / UNVERIFIABLE / STALE / FORWARD-LOOKING**. On-chain-checkable predicates today: supply fixed, no-mint, vesting parameters, LP locked (locker + expiry), renounced ownership, multisig threshold, timelock delay, allocation percentages, audit-hash-matches-bytecode.

- **Tense handling.** Present-tense claims ("liquidity IS locked") are verified now. Forward-looking claims ("LP WILL be locked at launch") become *tracked commitments* with deadlines — false only if the deadline passes

unmet. This is what turns trust into a time series.

- **Omissions count.** An on-chain power the docs never mention is charged as an omission. Honest self-disclosure of warts scores *up*.
- **LLMs never touch the math.** Extraction and explanation are LLM work; every verdict that affects the score requires a deterministic on-chain check or human review. This is what makes "unpurchasable algorithmic scores" defensible — and it is also the defamation firewall (Part VI).

2.3 Grade scale and aggregation

Scale: A+, A, A−, B+, B, B−, C, D, F, plus **N/R** (unratable) and a **Provisional** badge for young tokens. Numeric mapping: A ≥ 85, B 70–84, C 55–69, D 40–54, F < 40. Notches exist only at the top, where projects actually compete on quality. Every grade carries a **Trust Model badge** — *Code-Enforced / Governance-Gated / Custodial (disclosed) / Discretionary* — so a B+ stablecoin and a B+ fixed-supply token are visibly different kinds of B+.

Aggregation is a hybrid: `overall = min(weighted average, all triggered caps)`. A pure average fails the mint-backdoor test (perfect docs would buy a B despite an EOA mint). Pure weakest-link kills the gradient (fix one thing, grade doesn't move). The caps are the weakest links that are actually fatal; the average preserves the incentive to improve everywhere else.

The cap table (published verbatim as part of the methodology)

TRIGGER	OVERALL CAP
Unverified source code	C
Mint/issuance callable by an EOA or undisclosed party	F
Upgradeable with EOA admin, no timelock	D
Materially false safety claim (verified FALSE)	D
>50% of primary liquidity withdrawable by one key	D
Historical commitment violation	D (→ C after 24 clean months)
Insider liquid float >30%	C
Top-10 unlabeled holders >50%	C
Material-claim coverage <50% verifiable	B
Age <12 months	A− + Provisional
Age <6 months	B+ + Provisional

Each rating also shows a data-coverage indicator (● full / ◐ partial / ○ thin). A B+ at full coverage and a B+ at thin coverage must be distinguishable, or the first embarrassing miss destroys the brand. Category profiles (published) re-weight rather than fork the methodology: stablecoins are judged on purpose-consistent issuance and redeemability; claim-light memecoins get a disclosure floor of 70 but cap at B+ permanently; L1/L2 native coins (ETH, SOL) are out of automated scope in v1 — a few hand-written static profiles anchor the scale.

2.4 Sanity check — the famous-token test

TOKEN	DRIVERS	OUTCOME	BADGE
USDC	Mint-by-design with attestations; admin surface without timelock; excellent disclosure; long clean record	B+	Custodial (disclosed)
UNI	Capped governance-gated mint; Governor + Timelock; vesting long completed; deep organic liquidity	A-	Governance-Gated
SHIB	No mint, renounced; claim-light; whale concentration; multi-year clean record	B	Code-Enforced (claim-light)
Fresh pump.fun graduate	Structure fine by platform; heavy sniper float; zero docs; no history	D+/C- Provisional	Discretionary
DeFi token falsely claiming "LP locked"	Everything else A-grade — the false-claim cap binds	D	—

Every row matches educated-crypto intuition, including the deliberately spiky ones: SHIB above a lying DeFi token is a *feature* — it teaches the market what the grade means.

2.5 Gaming resistance

Master principle: **choose metrics where gaming them means actually becoming trustworthy**. If a project locks LP for real and vests for real to farm the score, Goodhart's law is working for us. The residual attacks and countermeasures:

ATTACK	COUNTERMEASURE
Lock roulette — lock LP 30 days, screenshot the badge	Score scales with duration; expiries monitored; grade auto-degrades on expiry (this is the paid alert product)
Fake locker — "lock" LP in a team-controlled contract	Allowlist of audited lockers; unknown locker = no credit
Renounce-but-retain — grant a minter role, then renounce owner	Scan full historical RoleGranted events, never just current owner()
Proxy sleight-of-hand (beacon/diamond patterns)	Detect all proxy patterns; historical Upgraded events; latent-mint treatment for live upgrade paths
Sybil insider float — fund 50 wallets pre-"fair launch"	Genesis forensics; unexplained early whales default to insider-attributed
Prompt injection in docs ("rate this token safe")	Docs are data, never instructions: extraction-only schema, quoted evidence spans, verdicts must cite chain evidence; adversarial injection suite in CI
Claim-stripping — delete all promises	Claim-light floor + coverage cap: silence can't grade excellent

Reputation reset — redeploy to erase a violation	Deployer-lineage flag; violations attach to team identity where provable
Buying the score	Structurally impossible: no issuer-pays; methodology and every grade change published with evidence; public appeals

2.6 Worked example: ATA Coin — the credibility exhibit

Running the methodology on the founder's own token (design audit; ATA is pre-deployment): Supply Integrity 100 (no mint path exists, atomic constructor distribution), Admin Surface 92 (no owner, no proxy; one disclosed nuance), Disclosure ~75 now → ~95 once the LP-lock transaction verifies (one material claim is still forward-looking), Insider Float 90 (85% code-vested, **insider liquid float at launch = 0**), Liquidity N/R until the lock exists, Track Record unrated pre-deployment.

Aggregate ≈ 91 (A- territory) — then the caps bind: age < 6 months → final grade B+ (Provisional), Code-Enforced. Path: B+ → A- at 12 months → A at 24 months, contingent on zero commitment violations. The rating system's own seed project cannot score an A on day one. Publish this worked example inside the methodology — it is the single best proof the scale isn't rigged.

Market & Positioning

Who's already here, which gaps are real, where the wedge goes in, and the honest kill risks.

3.1 The landscape in one table

PLAYER	WHAT IT IS	WHAT IT DOES NOT DO
GoPlus	Free checkbox API (honeypot, mintable, blacklist), ~717M calls/mo	No grade, no docs analysis, no time series — and only \$4.7M cumulative revenue ever
Token Sniffer (Solidus)	0–100 automated safety score, ~15 EVM chains	No docs/behavioral layer; explicitly "not an audit"
CertiK Skynet	Score across 17,000+ projects	Independence — revenue is issuer-paid audits (the positioning foil)
RugCheck	Free Solana checks (mint/freeze, LP burn, snipers)	Solana-only, checkbox, point-in-time — makes Solana a hard wedge
TokenInsight	AAA–D letters on ~247 tokens since 2018	Tiny coverage, slow, no automated verification, faded mindshare
L2Beat	Stage framework for rollups — the legibility model we borrow	L2s only; nonprofit; proves the format, not the business
Exponential.fi	A–F grades on DeFi pools; backtest: A-rated = zero user losses	Pools, not tokens; proves letter grades + open methodology work
Tokenomist	Unlock data + one "on-chain claims" feature on large caps	One claim type, no grade, no gap-scoring, no alerts
Hexagate / Hypernative	Enterprise behavioral monitoring (Hexagate → Chainalysis, ~\$60M)	Private B2B only — the public per-token record remains unbuilt

THE CERTIK FOIL — USE IT EXPLICITLY

Kraken (2024): CertiK researchers exploited a bug for ~\$3M and returned funds only after a public standoff. Merlin DEX (2023): rugged days after a CertiK audit via the very issue marked resolved. The structural criticism writes the positioning: **issuer-pays is the 2008 rating-agency conflict re-run; our scores cannot be bought.**

3.2 The wedge: Base/EVM tokens that make claims

Not pump.fun dust. The claims engine needs claims: pump.fun tokens have no website, no whitepaper, no vesting promise — nothing to verify — and graduation has collapsed to ~0.2%, with RugCheck already free and entrenched there. The Base/EVM mid-tail (project tokens, app-coins, agent-coins, launchpad and RWA/DeFi tokens) is where

projects *do* publish docs, lock LP, and promise vesting — exactly "serious enough to make promises, young enough to rug." It is also culturally adjacent to Coinbase's legibility ethos, and the founder's existing tooling is Base-native. Sequencing: win "every Base token with a website" (hundreds — tractable solo), then the EVM long tail, then Solana's graduated survivors.

3.3 Distribution, ranked

1. **X mention-bot** ("@bot grade \$TOKEN"). The loop is proven — aixbt went 0 → ~500k followers in about two months; Bankr drove \$4.7B volume from mention-driven UX. Cost reality under 2026 pay-per-use pricing (~\$0.005/read, \$0.015/post, **\$0.20 if the post contains a link**): a busy bot runs \$500–700/mo — affordable, but reply with the grade as an image/text and no hyperlink. Platform risk is real (X purged 1.7M bots in 2025–26), so:
2. **Farcaster / Base App** — free, open API, exactly the Base wedge audience.
3. **Telegram bot** — free API; traders already live in TG bots; grade-on-paste-contract.
4. **MCP server + x402 API for AI agents** — trading agents need a machine-readable trust signal before touching a token; nobody offers *grades* over MCP; the ATA MCP server ports directly. The cheapest "embed everywhere" play of 2026.
5. **Browser extension** overlaying grades on DEX Screener/Basescan — later, for retention.
6. **Wallet embeds** — deprioritize; wallets already bought their security layers (that's the acquisition path, not the launch path).

3.4 Kill risks, ranked honestly

1. **Monetization failure, not competition.** The graveyard says ratings are a feature, not a business (GoPlus: ~1B calls/mo, \$4.7M ever; L2Beat lives on grants; DefiSafety semi-dormant). Counter: the paid product is *alerts on trust changes* — urgent, personal, recurring — the one shape that demonstrably monetizes (Nansen alerts, \$49–69/mo).
2. **Incumbents add LLM claim-checks.** Counter: GoPlus's DNA is checkbox APIs, CertiK can't credibly grade its own paying clients — and the track-record time series cannot be backfilled. Speed matters; the window is open now.
3. **The no-docs problem.** Rug anxiety is hottest where there are no claims to verify. Counter: the behavioral grade still applies docs-free; don't promise memecoin coverage we can't differentiate.
4. **One viral wrong grade.** Counter: confidence tiers, human review before any low grade on a large token, public corrections log, versioned methodology.
5. **X platform risk.** Counter: multi-channel from day one; X is acquisition, the site is the asset.
6. **Legal threats over D/F grades.** Counter: the entire Part VI architecture.

3.5 Validation before building (weeks 1–4, ≈ \$0)

- **20-token manual pilot:** extract claims from 20 Base token sites with Claude, verify by hand/script, publish grades as a thread + one static page. Signal: "grade X next" replies, project responses, >50k organic impressions.
- **Mention-bot MVP (\$100–300):** grade-on-mention from the pilot's cached DB. Measure pull vs. aixbt's early curve.
- **Pre-sell 5–10 deep-dive reports (\$99–499 intro)** to launchpads, DAOs, small funds — tests the paid-report line before the infra exists.

- **Alert smoke test:** landing page + mock alert screenshots, \$15/mo waitlist with card capture. >100 signups = build it.
- **One embed conversation** with a Base launchpad or screener — even a "no" reveals the data shape they'd pay for.

Technical Architecture & Build Plan

Near-zero cost at MVP scale, a 6-week build, and the honest list of what cannot be automated.

4.1 What the ATA repository contributes

~300–500 directly liftable lines plus the architectural patterns. ATA reads a *known* deployment; the rating product must handle *arbitrary hostile tokens* — a different problem — but the plumbing transfers:

ASSET (IN ATA-COIN)	REUSE
<code>mcp/src/chain.ts</code> — fallback multi-endpoint RPC client	Lift directly: the free-RPC resilience layer
<code>site/js/rpc.js</code> — batched <code>eth_call</code> with <code>chainId</code> verification, zero deps	Lift directly: one round-trip per ~30 calls is the core cost lever
<code>site/js/abi.js</code> — selector table cross-checked in CI	Pattern: grow into a large selector registry (mint, blacklist, <code>setTax...</code>)
<code>mcp/src/tools.ts</code> — config-vs-chain cross-check	The docs-vs-chain truth engine in miniature — generalize claimed-vs-observed → verdict
<code>scripts/report.ts</code> — markdown report from chain state	Template for auto-generated per-token evidence reports
MCP server (7 read-only tools)	Port wholesale later as the trust-rating MCP server
GitHub Actions cron + static-site data export	The exact v1 shape: cron writes JSON → static site renders it

4.2 v1 shape

```

trustrating/
  packages/core      # verifier registry, types, scoring engine (pure functions, versioned)
  packages/chain     # rpc batch client (from ATA), selector + locker + label registries
  packages/claims    # crawl -> extract -> verify -> judge pipeline (Anthropic SDK)
  apps/scanner       # cron endpoint; writes SQLite + JSON
  apps/site          # static site reading built JSON (grades, evidence, methodology)
  apps/alerts        # (v2) watcher loop + Telegram/webhook fan-out
  data/              # git-committed registries, scan outputs, score history

```

- **Data sources, cheapest first:** public Base RPCs (batched `eth_call`; Alchemy free tier as backstop), Etherscan V2 API for verified source, DexScreener free API for pools/liquidity/social links (the crawl seed), Blockscout for holders, GoPlus as an advisory cross-check only, direct storage-slot reads for proxy detection, a curated locker registry (UNCX, Team Finance...) maintained in-repo.

- **Claims pipeline:** crawl (content-hashed, re-run only on change) → Haiku extracts to a strict schema → deterministic TypeScript verifiers → Sonnet judges and writes the rationale — all via the Batch API at ~50%. ≈ \$0.05–0.10 per token analyzed.
- **Storage:** SQLite, snapshot into git. Postgres only when the firehose demands it.
- **Deployment v1:** GitHub Actions cron + Cloudflare Pages, \$0. Static JSON files *are* the free API. Bonus: a public pipeline repo makes every score git-timestamped — nobody can accuse us of retroactively editing grades.
- **Monitoring (v2), two clocks:** an event clock (1–5 min eth_getLogs sweep over all watched addresses — a handful of calls covers thousands of tokens; Base's 2s blocks give ~2-minute alerts) and a calendar clock (unlock/cliff deadlines captured at scan time; fire T-30d/T-7d/T-24h and verify execution after). Severity: CRITICAL → instant fan-out + grade recompute; WARN; INFO. Every alert embeds evidence links — the alerts *are* the track record, so persist them all.

4.3 Cost model at three scales

SCALE	CHAIN READS	LLM	INFRA	TOTAL / MONTH
100 tokens (MVP)	Free public RPCs + free API tiers	\$10–20 one-time backlog; \$2–5/mo re-runs	GH Actions + CF Pages free	≈ \$5
10,000 tokens	getLogs deltas + weekly rescans: \$0–25	\$150–300 one-time; \$30–80/mo	Workers \$5 or small VPS	\$50–150
500,000 tokens (firehose)	Streaming ingestion; paid RPC tier \$50–300	Claims engine gated to ~\$50k+ liquidity tokens; chain-only provisional grades for the rest	Postgres \$25–70	\$400–1,200

The firehose scale is only entered when revenue carries it. LLM cost is real but never the constraint; daily behavioral monitoring is RPC reads and rounds to zero.

4.4 The 6-week MVP

Base only. 200–500 launch tokens (top Base tokens by liquidity + any requested). Six automated checks: supply/mint risk; upgradeability + who-controls; liquidity safety; holder concentration; honeypot/tax (advisory); and **docs-vs-chain lite** — five claim types (liquidity locked / renounced / vesting / fixed supply / audited) auto-verified and published as a per-token claims table. Even in reduced form, check #6 is the launch story: nobody else ships it.

WEEK	DELIVERABLE
1	Monorepo + chain package + checks 1–3 against 10 hand-picked tokens (golden-file tests; ATA as known-good fixture)
2	Checks 4–5 + scoring engine + methodology doc v0.1
3	Claims pipeline happy path
4	Site + 300-token batch scan + manual QA of top 50

5	Cron, rescan diffing, Telegram bot (own use)
6	Polish + launch content: five "trust-gap reports" on Base tokens where docs contradict chain

4.5 What cannot be automated (say so in the product)

1. **Off-chain claims** (team identity, partnerships, "backed by") surface as *unverifiable* — expect >50% of extracted claims here. Never verified, never contradicted.
2. **Audit claims are a trap:** score only "audit link resolves on the auditor's domain," nothing stronger. Fake-audit PDFs are common.
3. **Team-wallet identification is heuristic:** a false "team is dumping" alert is the worst failure mode legally and reputationally — require two independent signals before any CRITICAL wallet alert.
4. **Vesting compliance** only works for on-chain vesting; ATA-style transparency is the exception — which is the point of the product.
5. **LLM extraction failure modes:** puffery vs. claims, temporal ambiguity, stale docs, wrong-token confusion on multi-chain deployments — anchor every claim to the contract address the project itself publishes.
6. **Registry gaps and RPC failures must surface as "insufficient data"** (grade capped but labeled), never silently as the worst case.
7. **Human review is structurally required** for D/F grades on any token above ~\$1M liquidity, for paid reports, and for registry additions. Budget 2–4 founder-hours/week at MVP scale — the real bottleneck is QA, not compute.

Business Model & Financials

Free grades forever; charge for time-sensitivity, machine-readability, and depth.

Bootstrapped through month 18.

5.1 Revenue lines, ranked by 36-month potential

#	LINE	PRICE POINTS	36-MO POTENTIAL	NOTE
1	API licensing (wallets, screeners, bots, AI agents)	\$99–499/mo self-serve; \$1–2k+/mo enterprise	\$3–6k MRR base	GoPlus's ~717M calls/mo proves B2B demand; agents are the new buyer class nobody prices for
2	Pro subscriptions (behavioral alerts)	\$24/mo	\$8–12k MRR base	First revenue; price the time series, not the grade; churn is brutal (5–10%/mo) — modeled net
3	Deep-dive reports (buyer-pays)	\$2.5–7.5k each	\$30–50k/yr	Priced below the \$5k audit floor; funds, exchanges, DAOs, law firms
4	Grants (Base, Optimism Retro, Gitcoin)	\$10–100k+/yr	Front-loaded, years 1–2	L2Beat took ~\$600k in RetroPGF 3; open methodology + free data is exactly what qualifies
5	Embeds / attribution deals	\$0 early	Distribution, not revenue	Free API requires attribution — every integration advertises the brand

THE "EXPEDITED REVIEW, PUBLISHED EITHER WAY" IDEA — VERDICT

Structurally clean (the fee buys queue position, never the score), **poisonous in sequencing**: with no reputation, outsiders can't distinguish "paid for speed" from "paid for score," and the first expedited-paying project that rugs writes the headline. Do not offer it in year 1. If ever introduced (month 18+, after 500+ unpaid grades): flat public price, mandatory publication, fee disclosed on the report, capped at ~25% of revenue, no refunds for bad grades. If any guardrail feels commercially unbearable, that is the signal it has become issuer-pays in disguise — kill it.

5.2 Launch pricing

TIER	PRICE	CONTENTS
Free (web)	\$0	All grades, breakdowns, claim-level evidence, methodology, delayed timeline — free forever; the grades are the marketing
Free (API)	\$0	Grade + dimensions, ~30 req/min, attribution required ("Rated by ...")

Pro	\$24/mo	Real-time alerts (admin key, LP expiry, vesting breach), 50-token watchlist, Telegram/webhooks
API Growth	\$99/mo	10 req/s, claim-level detail, no attribution requirement
API Scale	\$399/mo	50 req/s, grade-change webhooks, batch, alert firehose
Enterprise	\$1–2k+/mo	SLA, historical dumps, white-label feed — quote on inbound only
Deep-dive report	\$3.5k list	Human-reviewed, AI-drafted trust-gap report; never issuer-pays at launch

One paid consumer tier only (Nansen's 2026 collapse to a single Pro tier is the market speaking); annual = 2 months free; grandfather early API integrators at \$0 for attribution + a testimonial.

5.3 Go-to-market (solo founder, zero audience)

- **Phase 0 — the credibility corpus (weeks 1–4):** publish the methodology and ~100 grades of well-known tokens, including famous rugs graded retroactively. **Grade ATA first, adversarially, and publish it** — "we graded our own founder's token; here is every gap" is a launch story only this founder can tell.
- **Phase 1 — the grading bot (months 1–6):** grade controversial trending tokens within hours of them trending; one damning verified claim ("site says liquidity locked 2 years; on-chain lock expires in 11 days") is quote-tweet fuel. Cadence: 2–3 proactive grade drops weekly + a weekly Trust-Gap Report thread + monthly grade migrations (trains the market that grades are alive).
- **Phase 2 — partnerships (months 4–12), in order:** ① Base ecosystem grants (Builder Grants are retroactive 1–5 ETH; Base Batches \$10k + \$50k investment track); ② small Base-native tools (trackers, TG bots, Farcaster mini-apps — they say yes in days); ③ DEX screeners (traffic + a differentiation problem; slow cycles, start early); ④ wallets last — pitch the complement: "GoPlus says it's a honeypot; we say the team's promises don't match the chain."

5.4 Financial sketch

	PESSIMISTIC	BASE	OPTIMISTIC
Month 12	~\$150 MRR + 1 report + 1 small grant → ~\$10k yr-1	~\$2.2k MRR + 4 reports + grants → ~\$60–80k yr-1	~\$10–15k MRR + \$30k reports + \$75k grants → ~\$200k yr-1
Month 36	~\$25k/yr	~\$15k MRR / ~\$180k ARR + \$40k reports → ~\$220–300k/yr	\$0.7–1.2M ARR — a fundable company
Cash break-even	Month 20+ or never	Month 8–12	Month 4–6

Cash burn: ~\$160–480/mo in months 1–6 (infra, LLM, X API, misc), ~\$500–1,400/mo in months 7–18. The pessimistic case is genuinely plausible — trust products have a cold-start paradox (nobody checks a rater with no reputation); the credibility corpus and grade-drop cadence exist to break it.

Stay bootstrapped through at least month 18. Costs are a few hundred dollars a month, independence is itself the product, and grants realistically cover the gap (\$15–50k year 1, \$30–100k year 2 if methodology and data stay genuinely open). Raise only for the embedded-standard end-state, and only with the leading indicators in hand (>1M free API calls/mo, 2+ named integrations, \$15k+ MRR) — then a \$1–2M round (Coinbase Ventures the obvious strategic) buys multi-chain coverage and a first hire.

5.5 Year 3–5 end-states — keep all three doors open

DOOR	SHAPE	REQUIRES
A — Lifestyle data business	\$200–500k/yr, founder + one contractor, ~85% gross margin	Not blowing up the cost structure; no growth-demanding capital
B — Acquisition	Realistic \$2–10M on base traction (comp: Token Sniffer → Solidus); \$10–30M with \$1M API ARR	Buyers: compliance platforms, security firms, exchanges/wallets, data aggregators — they want the longitudinal evidence dataset
C — The embedded standard	"The grade every frontend shows" — highest ceiling, most credibility-dependent	Free grades forever, open methodology, visible neutrality, eventually grant/consortium funding

Decisions made now that keep all doors open: ① no token for the rating service, ever; ② no issuer-pays for scores; ③ open the methodology and grade *data*, keep the claims-engine *code* proprietary — the single most important structural split; ④ store everything as structured evidence records from day 1 (the un-backfillable asset); ⑤ free API requires attribution; ⑥ clean single-member entity separate from ATA with all IP assigned; ⑦ no exclusive data deals.

Legal, Entity & the ATA Conflict

Publishing impersonal token ratings is legal everywhere that matters — as long as three properties hold. The architecture is the defense.

6.1 The three properties that keep it unregulated

Every regulatory trap in every jurisdiction is triggered by drifting from one of these: the product stays **impersonal** (same score shown to everyone), **unpaid-by-issuers**, and free of **personalized buy/sell output**. Alerts on objective events ("admin key changed") are fine; "you should sell" is not.

JURISDICTION	ANALYSIS
United States	Two independent shields: the 2026 SEC/CFTC framework classifies most fungible tokens as commodities (rating non-securities isn't investment advice), and the publisher's exclusion (<i>Lowe v. SEC</i>) covers impersonal publications of general circulation regardless. NRSRO rules are scoped to <i>credit</i> ratings — confirmed non-trigger. Branding rule: never call it a "credit rating," don't imitate AAA/Aaa trade dress.
European Union	MiCA's only analysis-shaped licensable service is <i>personalised</i> advice — a public identical-for-everyone rating is not a CASP service. EU CRA Regulation: debt-scoped, non-trigger. Watch item: the ESG Ratings Regulation (applicable July 2026) brushes "governance scores" — mitigate by framing dimensions as protocol/technical integrity, and revisit before signing any EU-regulated financial institution as an API customer.
United Kingdom	The trap is the cryptoasset financial-promotion regime (criminal since Oct 2023). Neutral ratings are not inducements. Design rules: no exchange affiliate links, no "buy here" buttons, no "top tokens to buy" framing, no sponsored placements — which is what brand neutrality wanted anyway.
Türkiye (founder's residence)	A ratings publisher is not a CASP under Law 7518, but Türkiye criminalizes unauthorized investment advisory (3–5 years) and prosecutes paid "signal groups." The paid tier must sell monitoring of objective events, never buy/sell calls. Operate through the foreign entity, English-language, no Türkiye-targeted marketing, standard "yatırım tavsiyesi değildir" disclaimer. Get 1–2 hours of Turkish capital-markets counsel before launch (~\$300–600) and a Turkish CPA opinion on LLC income characterization at first revenue (~\$500–1,500).

6.2 Defamation — the #1 real risk, and the product architecture as the defense

US doctrine: **an opinion based on fully disclosed, true facts cannot be defamatory, no matter how harsh**. What is actionable: false factual assertions ("the team dumped") — and disclaimers do not immunize them. The instructive precedents: rating agencies lost First-Amendment cover when ratings went to a select group (free public ratings get the strongest protection — keep the grade and its factual basis public even when analytics are paid), and short-sellers lose when specific factual claims are false and a profit motive supplies malice (hence the no-trading rule).

Launch-blocking design requirements — not legal nice-to-haves:

1. **Fact/opinion separation in the data model:** every dimension = (verifiable fact, evidence link, timestamp, method) + (grade = labeled opinion derived by published formula). The UI shows the chain: claim → check → result → grade.

2. **Evidence snapshots** behind every negative finding (tx hash, contract state, docs capture with date) — the truth file.
3. **LLM output verification gate:** a hallucinated "team wallet dumped 40%" is a false factual assertion published at scale. No extracted claim or verdict publishes without a deterministic on-chain check or human review.
4. **Calibrated language:** "no lock found at the stated address as of block N," never "they lied."
5. **Dispute & correction process** with a 10-business-day SLA, wallet-signature authentication of disputants, public changelog, and one linked official response per rated project.
6. **Forum choice:** US LLC + terms selecting a strong anti-SLAPP state (California, Texas, Nevada) for early dismissal and fee-shifting.
7. **Insurance:** media-liability/E&O (~\$2–6k/yr) at first revenue or first legal threat, whichever comes first.

6.3 Entity: Wyoming LLC

OPTION	YEAR 1	RECURRING	STRIPE?	VERDICT
Wyoming LLC	~\$160–300	~\$450–700/yr	Yes	Do this — before the first published F grade
TR sole proprietorship	~\$200	+CPA monthly	No (Stripe unavailable in TR)	No liability shield for defamation
Delaware LLC	~\$510–710	~\$400+/yr	Yes	Only if VC path; pricier for nothing needed now
Estonia OÜ / UAE freezone	€2.2k / ~\$4k	€1.2–1.8k / \$6–9.5k	Yes	4–10× the cost; revisit UAE only at six-figure profit

Sequence: form the LLC + EIN (free, 2–6 weeks without an SSN) → banking via Mercury (have a real business narrative ready; 2026 Mercury rejects registered-agent-only addresses) or Airwallex/Wise/Relay → payments via Stripe, or simpler on day one: a merchant of record (Paddle / Lemon Squeezy / Polar) which works for Türkiye-based sellers and handles global VAT. US tax: foreign-owned single-member LLC with no US office ⇒ \$0 US income tax, but **Form 5472 + pro-forma 1120 annually — the one deadline that matters (\$25k penalty)**. Türkiye taxes worldwide income; the LLC-characterization question is the one paid opinion to buy.

6.4 Data licensing — the hidden landmine

Etherscan's API terms prohibit commercial use and AI/dataset use. Do not build the product on it. Build on: public JSON-RPC endpoints (reading chain state is permissionless), Blockscout (open source), Alchemy/Infura free tiers (commercial use allowed), The Graph. The existing ATA chain-reading code on public RPCs is already the compliant pattern. Scraping project docs: public pages only, never click-accept a ToS, respect robots.txt, stop per-site on a cease-and-desist (and note "project blocks verification" — itself signal). Quote short claim excerpts (fair use for verification/criticism); keep full-page snapshots private as litigation evidence.

6.5 The ATA conflict of interest — strictest version, published

1. **ATA's page shows "Not Rated — Founder Conflict of Interest"** with the raw dimension facts and no grade (cleanest), or a purely algorithmic grade behind a permanent banner. Never quietly rate it well.
2. **Site-wide disclosure:** the methodology page lists the founder's token holdings and affiliations, updated on change.
3. **Competitor tokens:** algorithmic output only — any human override on a competitor's grade is prohibited; their pages carry an automatic disclosure.
4. **No-trading rule, published:** no trading any rated token around its rating actions; never short rated tokens.
5. **Origin story disclosed:** "the methodology grew out of the adversarial trust review built for ATA" is a credibility asset if disclosed, a scandal if discovered.

6.6 Priority legal actions

#	ACTION	COST	WHEN
1	Form Wyoming LLC + EIN	~\$160–300	Before first published rating
2	Fact/opinion/evidence data model written into the product spec	\$0	Now
3	ToS + disclaimers + dispute/correction + COI policy (AI-drafted, 2h attorney review)	~\$300–600	Launch
4	Turkish CPA opinion (LLC characterization, CFC)	~\$500–1,500	First revenue
5	TR capital-markets counsel sanity check	~\$300–600	Launch
6	MoR (Paddle/Lemon Squeezy) or Stripe + Mercury under the LLC	fees only	First paid tier
7	Zero Etherscan-API dependency in the build	\$0	Build phase
8	Media-liability insurance; trademark filing (name + badge)	\$2–6k/yr; ~\$350	Post-revenue / first threat

Total year-one legal/entity spend $\approx$ \$1,500–3,000. Nothing here blocks launch. The two things that genuinely hurt later: a hallucinated factual claim about a litigious project (solved by the verification gate, not disclaimers), and drift toward personalization or issuer monetization (which converts an unregulated publisher into a regulable adviser everywhere at once).

Name & Domains: What to Buy

The direct answer to "tell me what domain to buy" — a ranked recommendation with an explicit fallback trigger.

7.1 Recommendation: Assay

"**Assay**" is already a verb for exactly this product — the centuries-old word for testing a precious metal to determine its true content. "We assayed the token" needs no explanation; the metaphor (marketing says 24-karat, the assay says what's really in the bar) is the product. It is short, professional, pronounceable, works as "an Assay grade," resonates with Turkish *ayar* (the fineness/karat rating of gold — a private but pleasing echo for the founder), and as an arbitrary-ish mark applied to crypto ratings it is trademark-strong. Naming research screened the crypto/fintech space and found no active collision in this category.

THE PURCHASE ORDER — DO ALL OF IT IN THE SAME HOUR

1. **assay.xyz** — primary candidate. ⚠ Single-word .xyz domains are often registry-premium; check the renewal price, not just year 1.
2. **assay.io** — secondary. Same premium caution.
3. **getassay.com** and **assayratings.com** — defensive + email/redirect. Should be standard-price.
4. **assay.eth** — ENS name (crypto-native identity; cheap, annual).
5. **X handle** (@assay / @assayratings / @assayxyz — best available) and **GitHub organization** — register in the same hour as the domains; handles leak.

Budget: \$50–150/year in the normal case; a one-time aftermarket purchase is acceptable up to ≈ \$1,500 if assay.xyz or assay.io is buy-now-priced. Do not pay a premium *renewal* above ~\$500/yr — recurring premiums are a tax forever.

FALLBACK TRIGGER — DECIDE AT THE REGISTRAR, SAME DAY

If **assay.xyz** and **assay.io** are both taken or premium-renewal above ~\$500/yr, fall back to **TokenFacts** and buy **tokenfacts.com** + **tokenfacts.xyz** + **tokenfacts.io** the same day (tokenfacts.com appeared inactive at research time). TokenFacts is the descriptive, does-what-it-says alternative — weaker as a trademark, instantly clear as a product. Runners-up if both fail: TokenLabel, TokenBeat, OnchainFacts. Names screened and eliminated for active collisions: TrustGrade, ChainFacts, TokenGrade, ClaimCheck, Touchstone, and the Veri-/Vera- family.

Honest caveat: live registrar availability could not be verified from this research environment (domain-lookup services are blocked here); the collision screening was done via web research. Availability and premium status must be checked at the registrar (Namecheap, Porkbun, Cloudflare Registrar) at the moment of purchase — hence the same-day fallback trigger above. Turn on WHOIS privacy and auto-renew for everything; put the domains in the LLC's name once it exists.

7.2 Brand hygiene from day one

- Same name for the site, X/Farcaster/Telegram handles, GitHub org, and MCP server — the anti-impersonation lesson from ATA generalizes.
- File the US trademark (~\$350/class) once revenue justifies it; the "Rated A by Assay" badge is the thing the trademark protects (badge license: must fetch the live grade, no stale or selective display).
- Never brand anything a "credit rating" (Part VI branding rule).

PART VIII

Roadmap to End-State

Five phases, each with a gate. Nothing advances on hope; each phase must earn the next.

Weeks 1–4 Phase 0 — Validate before building

Spend ≈ \$0 and four weeks to find out if anyone cares, while the name and entity paperwork run in parallel.

- ☐ Buy the domains + handles (Part VII) — same hour.
- ☐ Start Wyoming LLC formation + EIN (slowest item — 2–6 weeks; start now).
- ☐ Run the 20-token manual pilot; publish grades as a thread + one static page.
- ☐ Alert-product smoke test (\$15/mo waitlist with card capture).
- ☐ Pre-sell 2–3 deep-dive reports at intro pricing.
- ◆ **Gate:** real pull exists (grade requests, >50k impressions, waitlist >100, or a pre-sold report). If nothing bites after an honest push, stop here — total sunk cost ≈ \$300.

Weeks 5–10 Phase 1 — MVP build & launch

The 6-week build (Part IV): Base only, 200–500 tokens, six checks, docs-vs-chain lite, static site, methodology v0.1.

- ☐ Ship the credibility corpus: ~100 well-known grades incl. retroactive famous rugs + the adversarial ATA self-grade.
- ☐ Launch content: five trust-gap reports on Base tokens where docs contradict chain.
- ☐ Mention bots live on Farcaster + Telegram; X bot linkless.
- ◆ **Gate:** organic grade-requests arriving weekly; one semi-viral grade drop; zero published factual errors (corrections log may have entries — that's fine, silence isn't).

Months 3–6 Phase 2 — Monitoring & first revenue

The two-clock watcher, alert taxonomy, paid Pro tier via MoR/Stripe, grade-migration cadence.

- ☐ Pro alerts at \$24/mo; free tier = 3 watched tokens.
- ☐ Apply: Base Builder Grants (retroactive — ship first), Gitcoin round, Optimism Retro.
- ☐ First integrations: small Base-native tools, free-with-attribution.
- ◆ **Gate:** first 50 paying subscribers or one grant landed; alert precision holding (no false CRITICAL on team wallets).

Months 6–12 Phase 3 — Moat build-out

The assets competitors cannot backfill.

- ☐ Track-record database compounding daily (started logging in Phase 1 — every day is unrecoverable advantage).
- ☐ Labeled claims corpus (claims → verdicts → outcomes); locker + label registries as public goods.
- ☐ Trust-rating MCP server + x402 API for AI agents; API tiers live.
- ☐ Add Ethereum + one more EVM L2 (same code path; resist Solana — it's a second stack).
- ◆ **Gate:** ~\$2k MRR or a screener/wallet conversation converting; break-even in sight (base case: months 8–12).

Years 2–5 Phase 4 — Choose the door

Doors A/B/C (Part V §5.5) stay open if the seven structural decisions held. Choose deliberately around month 18–30.

- ☐ Lifestyle: hold costs, deepen coverage, one contractor.
- ☐ Acquisition: keep evidence dataset pristine; entertain compliance/security/exchange buyers only with a real track record.
- ☐ Standard: raise the \$1–2M only with >1M free API calls/mo, 2+ named integrations, \$15k+ MRR.

Fusion with ATA: the two ventures stay separate, but each strengthens the other — ATA is the rating agency's best worked example, and a strong public grade earned over 24 clean months is the most credible marketing ATA could ever buy. Never let the money or the entities mix.

New Repository Bootstrap & Decision Log

What moves into the standalone repo, and every decision this plan has already made.

9.1 Repository bootstrap (when you say go)

1. **New GitHub org** (the name from Part VII) → new private repo `engine` (the pipeline; stays proprietary) + public repo `methodology` (the spec, registries, and published grade data — the grant-eligible public good). The public/private split mirrors decision ③ below.
2. **Copy from ATA-Coin** (copy, don't fork — separate history, separate license headers): `mcp/src/chain.ts`, the `site/js/rpc.js` batching pattern, the selector-registry + CI cross-check pattern from `site/js/abi.js`, the report-generator shape from `scripts/report.ts`, and the hardened CI workflow patterns (SHA-pinned actions, least-privilege tokens, `--ignore-scripts`).
3. **First-week commits**: monorepo skeleton (Part IV layout) → chain package with golden-file tests against 10 hand-picked Base tokens (ATA as the known-good fixture) → scoring engine as pure functions with `methodology_version` stamped on every result → methodology v0.1 in the public repo.
4. **Six research reports** from this planning cycle get committed to the new repo's `docs/research/` as founding documents (they are in the session scratchpad; ask and they'll be packaged).

9.2 Decision log — settled by this plan

#	DECISION	WHY
1	Product = algorithmic letter grades + docs-vs-chain claims engine + behavioral time series	The two verifiably empty niches; everything else is commoditized
2	Name: Assay; fallback TokenFacts on the registrar trigger	Part VII
3	Open methodology + grade data; proprietary pipeline code	Grant-eligible public good and acquirable IP at once — the load-bearing structural split
4	No issuer-pays for scores; expedited-review deferred to month 18+ or never	The CertiK foil is the positioning; don't become it
5	No token for the rating service, ever	Forecloses acquisition, taints neutrality — doubly so given ATA
6	Wedge: Base/EVM tokens with claims; not pump.fun; Solana later if ever	The claims engine needs claims; founder tooling is Base-native
7	Wyoming LLC before first published rating; MoR payments day one	Defamation shield + solves Stripe-for-Türkiye at minimum cost
8	Paid tier = event monitoring, never buy/sell calls	The Türkiye advisory line and the US publisher's exclusion, simultaneously

9	ATA page: Not Rated — Founder COI (or algorithmic-only + banner); published no-trading rule	Strictest COI version is cheap and load-bearing
10	No Etherscan API in the build; public RPC + Blockscout + licensed tiers	Etherscan ToS prohibits commercial/AI use
11	LLMs extract and explain; deterministic code decides; humans gate D/F on large tokens	Simultaneously the scoring integrity, the defamation firewall, and the injection defense
12	Validate (Phase 0) before building; every phase has a gate	The graveyard is full of unmonetized raters; \$300 buys the answer

9.3 What the founder does next (in order)

- ☐ Check assay.xyz / assay.io at a registrar → execute the Part VII purchase order (or the TokenFacts fallback) — same hour, all items.
- ☐ Start the Wyoming LLC + EIN paperwork (longest lead time).
- ☐ Say the word, and the new repo gets bootstrapped per \$9.1 and Phase 0 begins: the 20-token pilot is the first deliverable.

Synthesized from six commissioned research workstreams (methodology · market · technical · business · legal · naming), July 2026. Full reports available on request; sources cited within each. This document is a plan, not legal, tax, or investment advice — the flagged items in Part VI need paid professional opinions at the stated milestones.